

INFO 6105

Data Science Engineering Methods and Tools

Lecture 2

Linear Regression: A Probabilistic Approach

Ebrahim Nasrabadi
nasrabadi@northeastern.edu

College of Engineering
Northeastern University

Fall 2019

Recall from last lecture

Training examples:

$$(\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots, (\vec{x}_n, y_n)$$

where $\vec{x}_i = (x_{i1}, \dots, x_{im})^T \in \mathbb{R}^m$ is the feature vector and $y_i \in \mathbb{R}$ is the target value for $i^{th} example$.

We seek a function/hypothesis $h : \mathbb{R}^m \rightarrow \mathbb{R}$ such that $h(x)$ is a good predictor for corresponding values of y .

Recall from last lecture

Training examples:

$$(\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots, (\vec{x}_n, y_n)$$

where $\vec{x}_i = (x_{i1}, \dots, x_{im})^T \in \mathbb{R}^m$ is the feature vector and $y_i \in \mathbb{R}$ is the target value for i^{th} example.

We seek a function/hypothesis $h : \mathbb{R}^m \rightarrow \mathbb{R}$ such that $h(x)$ is a good predictor for corresponding values of y .

Modeling Choice: We assume that dependency of y on \vec{x} is linear and approximate y as a linear function of x :

$$h_{\beta}(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_m x_m = \vec{\beta}^T \vec{x}$$

where

$$\vec{\beta} = (\beta_0, \beta_1, \dots, \beta_m)^T \in \mathbb{R}^{m+1}$$

$$\vec{x} = (1, x_1, \dots, x_m)^T \in \mathbb{R}^{m+1}$$

How to determine β ?

Find $\vec{\beta}$ such that the linear model fits the training examples well

input	actual	predicted	residual/error
x_1	y_1	$\hat{y}_1 = \vec{\beta}^T \vec{x}_1$	$e_1 = y_1 - \bar{y}_1$
\vdots	\vdots	\vdots	\vdots
x_i	y_i	$\hat{y}_i = \vec{\beta}^T \vec{x}_i$	$e_i = y_i - \bar{y}_i$
\vdots	\vdots	\vdots	\vdots
x_n	y_n	$\hat{y}_n = \vec{\beta}^T \vec{x}_n$	$e_n = y_n - \bar{y}_n$

How to determine β ?

We define Residual Sum of Squares (RSS) and Mean Square Error (MSE) as

$$\begin{aligned}\text{RSS} &:= e_1^2 + e_2^2 + \dots + e_n^2 = \sum_{i=1}^n e_i^2 \\ &= \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2 \\ \text{MSE} &:= \text{RSS}/n\end{aligned}$$

This is sometimes called the loss function and denoted by $L(\beta)$.

How to determine β ?

We define Residual Sum of Squares (RSS) and Mean Square Error (MSE) as

$$\begin{aligned}\text{RSS} &:= e_1^2 + e_2^2 + \dots + e_n^2 = \sum_{i=1}^n e_i^2 \\ &= \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2 \\ \text{MSE} &:= \text{RSS}/n\end{aligned}$$

This is sometimes called the loss function and denoted by $L(\beta)$.
Choose β so as to minimize RSS.

Ordinary Least Squares (OLS)

Simple Linear regression

$$h(x) = \beta_0 + \beta_1 x$$

$$\frac{\partial L}{\partial \beta_0} = 0 \quad \frac{\partial L}{\partial \beta_1} = 0$$

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

$$= \frac{\text{Cov}(X, Y)}{\text{Var}(X)}$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

Ordinary Least Squares (OLS)

Simple Linear regression

$$h(x) = \beta_0 + \beta_1 x$$

$$\frac{\partial L}{\partial \beta_0} = 0 \quad \frac{\partial L}{\partial \beta_1} = 0$$

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

$$= \frac{\text{Cov}(X, Y)}{\text{Var}(X)}$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

Multiple Linear regression

$$h(\vec{x}) = X\vec{\beta}$$

$$\begin{aligned} L(\vec{\beta}) &= \|\vec{y} - X\vec{\beta}\|^2 \\ &= (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta}) \\ &= \vec{y}^T \vec{y} - 2\vec{y}^T X\vec{\beta} + \vec{\beta}^T X^T X \vec{\beta} \end{aligned}$$

$$\frac{\partial L}{\partial \vec{\beta}} = 0 - 2\vec{X}^T \vec{y} + 2\vec{X}^T X \vec{\beta}$$

$$\frac{\partial L}{\partial \vec{\beta}} = 0 \implies \vec{\beta} = (\vec{X}^T X)^{-1} \vec{X}^T \vec{y}$$

This lecture

Last time

- Discussed OLS as minimizing sum of square errors
- Gave explicit formula to compute parameters for the minimum

This lecture

Last time

- Discussed OLS as minimizing sum of square errors
- Gave explicit formula to compute parameters for the minimum

This time

- We rephrase the goals of OLS in a probabilistic light
 - ▶ Quick review of probability
 - ▶ Probability distributions
 - ▶ Likelihood function
 - ▶ Maximum likelihood estimation
 - ▶ Likelihood function and maximum likelihood estimates for linear regression

Quick review of probability

- Random Experiments
- Outcomes
- Probabilities
- Random Variables
- Expectation, Variance of a random variable
- Continuous random variables
- Probability density

Basic Definitions

Random Experiment: An experiment whose outcome is uncertain/non-deterministic.

Outcome : A realization of the random experiment

Sample Space Ω : Set of all possible outcomes of the random experiment

Event E : A subset of Ω is an event

Basic Definitions

Random Experiment: An experiment whose outcome is uncertain/non-deterministic.

Outcome : A realization of the random experiment

Sample Space Ω : Set of all possible outcomes of the random experiment

Event E : A subset of Ω is an event

Random Experiment: Roll a Die

Outcomes: 1,2,3,4,5,6

$\Omega = \{1, 2, 3, 4, 5, 6\}$

Events: $\{1\}, \{1, 2\}, \{1, 3, 5\}, \dots$



Random Experiment: Throw a dart at a dart-board

Outcomes: Any point (x, y) on the dart-board

$\Omega = \{(x, y) : (x, y) \in \text{the dart-board}\}$

Events: $E = \{(x, y) : (x, y) \in \text{inner circle}\}, \dots$



Probability of an Event

Probability: A mathematical quantity that defines how likely it is for an event to occur in a random experiment.

Probability of an Event

Probability: A mathematical quantity that defines how likely it is for an event to occur in a random experiment.

Mathematically,

$$P : \text{Set of events} \longrightarrow [0, 1]$$

that satisfies the following conditions

Laws of probability

(1) If A_1, A_2, \dots are disjoint events then

$$P(A_1 \cup A_2 \cup \dots) = P(A_1) + P(A_2) + \dots$$

(2)

$$P(\Omega) = 1$$

Interpretation of Probability

Two Interpretations:

Frequentist: Probability equals the long-run frequency of the event to occur, if the experiment was repeated several times.

Examples:

- The long-run frequency of heads in repeated tosses of a fair coin is $1/2$.
- The long-run frequency of 1's in repeated rolling of a fair die is $1/6$.

Subjective: Probability measures degree of subjective belief about uncertainty.

Examples:

- Probability that project X fails is 0.2.

Random variable: a variable whose value depends on possible outcomes of the randomized experiment.

Mathematically, is a mapping from the sample space Ω to a well-defined set.

- Real-valued $X : \Omega \longrightarrow \mathbb{R}$
- Discrete-valued $X : \Omega \longrightarrow \mathbb{N}$
- Binary-valued $X : \Omega \longrightarrow \{0, 1\}$
- Categorical $X : \Omega \longrightarrow \{\text{Excellent, Very Good, Good, Fair, Poor}\}$

Random Variable

Example: Suppose you throw a coin.

Sample space $\Omega = \{H, T\}$ Define a random variable:

$$X = \begin{cases} 1 & \text{if head occurs} \\ 0 & \text{if tail occurs} \end{cases}$$

For a fair coin,

$$P(X = 1) = P(X = 0) = 0.5$$

otherwise,

$$p(X = 0) = p \quad p(X = 1) = 1 - p$$

where

$$p = \frac{\text{number of times you observe a head}}{\text{total number of throws}}$$

will converge to p as the number of throws increases.



Random Variable

Random variables can be discrete or continuous.

- **Discrete** random variables have a countable number of outcomes

Examples:

- ▶
- ▶ {Excellent, Very Good, Good, Fair, Poor},
- ▶ number of received calls
- ▶ click or non-click

- **Continuous** random variables have an infinite continuum of possible values Examples:

- ▶ weight, volume of a package
- ▶ Time between calls to the customer service

Probability mass function

A **probability mass function (pmf)** is a function that gives the probability that a discrete random variable is exactly equal to some value.

Suppose that $X : \Omega \longrightarrow A$ is a discrete random variable defined on a sample space Ω . Then the probability mass function $f_X : A \longrightarrow [0, 1]$ for X is defined as:

$$f_X(x) = Pr(X = x) = Pr(\{s \in S : X(s) = x\})$$


Note:

$$\sum_{x \in A} f_X(x) = 1$$

Probability mass function



Probability Distribution Function (pdf)



x	$p(x)$
1	$p(x=1)=1/6$
2	$p(x=2)=1/6$
3	$p(x=3)=1/6$
4	$p(x=4)=1/6$
5	$p(x=5)=1/6$
6	$p(x=6)=1/6$

1.0

Probability density function


A **probability density function (pdf)** of a continuous random variable is a function whose value at any given sample (or point) in the sample space can be interpreted as providing a relative likelihood that the value of the random variable would equal that sample.

The probability density function ("p.d.f.") of a continuous random variable X with sample space Ω is an integrable function $f(x)$ satisfying the following:

- 1 $f(x) \geq 0$ for all $x \in \Omega$
- 2 $\int_{\Omega} f(x) dx = 1$
- 3 $P(a \leq X \leq b) = \int_a^b f(x) dx$

The Exponential Distribution

Generally the exponential distribution describes waiting time between Poisson occurrences.



$$P(1 \leq x \leq 2) = \int_1^2 e^{-x} = -e^{-x} \Big|_1^2 = -e^{-2} - (-e^{-1}) = -.135 + .368 = .23$$

$$P(1 \leq x \leq 2) = P(x \leq 2) - P(x \leq 1) = F(2) - F(1) = 2.3$$

Expectation

The **expected value** (or mean) of a random variable, intuitively, is the long-run average value of repetitions of the experiment it represents.

- Discrete random variable X ,

$$E(X) = \sum_i^n x_i P(x_i)$$

- Continuous random variable X ,

$$E(X) = \int_{\Omega} x f(x)$$

- Example: 6 faces fair dice, r.v. X is the value of one experiment

$$E(X) = \sum_{i=1}^n x_i P(x_i) = \frac{1}{6}(1 + 2 + 3 + \cdots + 6) = 3.5$$

- Example: exponential distribution X , $f(x) = e^{-x}$ ($x > 0$)

$$E(X) = \int x f(x) dx = \int_0^{\infty} x e^{-x} = 1$$

Variance is the expectation of the squared deviation of a random variable from its mean, and it informally measures how far a set of (random) numbers are spread out from their mean.

$$\text{Var}(X) = E[(X - E(X))^2] = E(X^2) - E^2(X)$$

- Example: 6 faces fair dice, r.v. X is the value of one experiment

$$E(X^2) = \frac{1}{6}(1^2 + 2^2 + \dots + 6^2) = 15.16, E(X) = 3.5$$

$$Var(X) = E(X^2) - E^2(X) = 15.16 - 3.5^2 = 2.92$$

- Example: exponential distribution

$$E(X^2) = \int_0^{\infty} x^2 f(x) dx = \int_0^{\infty} x^2 e^{-x} dx = 2$$


$$Var(X) = E(X^2) - E^2(X) = 2 - 1 = 1$$

Normal distribution



The probability density function of a Gaussian distribution is given by

$$p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(x-\mu)^2}$$

Where μ is the mean and σ^2 is the variance.



Standard Normal distribution



Question Why might linear regression, and specifically why might the least-squares cost function be a reasonable choice?

We give a set of probabilistic assumptions, under which least-squares regression is derived as a very natural algorithm.

Linear Regression Assumptions

Given training examples:

$$(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$$

we assume that $y_i = f(x_i) + \epsilon_i$ where

- $f(x_i) = \beta_0 + \beta_1 x_i$ (linear assumption)
- ϵ_i is a random term that captures either unmodeled effects, or random noise
 - ▶ $\epsilon_1, \dots, \epsilon_n$ are distributed IID (independently and identically distributed) according to a Gaussian distribution mean zero and some variance σ^2 , i.e.,

$$\epsilon_i \sim N(0, \sigma^2)$$

Recall that the density of ϵ_i is given by

$$p(\epsilon_i) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{\epsilon_i^2}{\sigma^2}\right)$$

Linear Regression Assumptions

In other words, we assume that each label y_i is Gaussian distributed with mean $\vec{\beta}^T \vec{x}_i$ and variance σ :

$$y_i \sim N(\vec{\beta}^T \vec{x}_i, \sigma^2)$$

This implies that

$$p(y_i | \vec{x}_i; \vec{\beta}) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - \vec{\beta}^T \vec{x}_i)^2}{2\sigma^2}\right)$$

The notation “ $p(y_i | \vec{x}_i; \vec{\beta})$ ” indicates that this is the distribution of given \vec{x}_i and parameterized by $\vec{\beta}$

Sales vs Price

Price	Avg. Sales
\$10.99	5.89
\$11.99	5.29
\$12.99	4.76
\$13.99	4.61
\$14.99	3.82
\$15.99	3.34
\$16.99	3.35
\$17.99	3.21
\$18.99	3.08
\$19.99	3.01
\$21.99	2.93

Sales vs Price

Given price, sales is a random variable with normal distribution parameterized by β .

Likelihood: probability of data given the parameters of the distribution
The probability of the data is given by

$$\begin{aligned} p(\vec{y}|X; \vec{\beta}) &= p((y_1, \dots, y_n) | \vec{X}; \vec{\beta}) \\ &= \prod_{i=1}^n p(y_i | \vec{x}_i; \vec{\beta}) \end{aligned}$$

This quantity is typically viewed a function of (and perhaps X), for a fixed value of β

When we wish to explicitly view this as a function of β , we will instead call it the **likelihood** function:

$$L(\beta) := L(\beta; X, \vec{y}) = p(\vec{y} | \vec{X}; \vec{\beta})$$

Maximum Likelihood

Maximum Likelihood Estimation: Find the parameters $\vec{\beta}$ such that, under normal distribution parameterized by $\vec{\beta}$, the observed data is most likely to occur. i.e., we should choose $\vec{\beta}$ to maximize $L(\vec{\beta})$

$$\begin{aligned} L(\vec{\beta}) &:= p(\vec{y}|X; \vec{\beta}) = p((y_1, \dots, y_n) | \vec{X}; \vec{\beta}) \\ &= \prod_{i=1}^n p(y_i | \vec{x}_i; \vec{\beta}) \\ &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - \vec{\beta}^T \vec{x}_i)^2}{2\sigma^2}\right) \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma}\right)^n \exp\left(-\frac{\sum_{i=1}^n (y_i - \vec{\beta}^T \vec{x}_i)^2}{2\sigma^2}\right) \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma}\right)^n \exp\left(-\frac{(\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta})}{2\sigma^2}\right) \end{aligned}$$

Maximum Likelihood

Maximum Likelihood Estimation: choose $\vec{\beta}$ to maximize $L(\vec{\beta})$. Instead of maximizing $L(\vec{\beta})$, we can also maximize any strictly increasing function of $L(\vec{\beta})$.

The derivations will be a bit simpler if we instead maximize the **log likelihood** $\ell(\vec{\beta})$

$$\begin{aligned}\log L(\vec{\beta}) &:= \log \left(\left(\frac{1}{\sqrt{2\pi}\sigma} \right)^n \exp \left(-\frac{(\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta})}{\sigma^2} \right) \right) \\ &= \log \left(\frac{1}{\sqrt{2\pi}\sigma} \right)^n + \log \left(\exp \left(-\frac{(\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta})}{\sigma^2} \right) \right) \\ &= n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{2\sigma^2} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta})\end{aligned}$$

Maximum Likelihood

Maximum Likelihood Estimation: choose $\vec{\beta}$ to maximize $L(\vec{\beta})$ or equivalently $\log L(\vec{\beta})$.

$$\log L(\vec{\beta}) = n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{2\sigma^2} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta})$$

Maximum Likelihood

Maximum Likelihood Estimation: choose $\vec{\beta}$ to maximize $L(\vec{\beta})$ or equivalently $\log L(\vec{\beta})$.

$$\log L(\vec{\beta}) = n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{2\sigma^2} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta})$$

$$\begin{aligned} \frac{\partial \log L(\vec{\beta})}{\partial \vec{\beta}} &= 0 - \frac{1}{2\sigma^2} \frac{\partial}{\partial \vec{\beta}} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta}) \\ &= \frac{1}{2\sigma^2} \frac{\partial}{\partial \vec{\beta}} \left(-\vec{X}^T \vec{y} + \vec{X}^T X \vec{\beta} \right) \end{aligned}$$

Maximum Likelihood

Maximum Likelihood Estimation: choose $\vec{\beta}$ to maximize $L(\vec{\beta})$ or equivalently $\log L(\vec{\beta})$.

$$\log L(\vec{\beta}) = n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{2\sigma^2} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta})$$

$$\begin{aligned} \frac{\partial \log L(\vec{\beta})}{\partial \vec{\beta}} &= 0 - \frac{1}{2\sigma^2} \frac{\partial}{\partial \vec{\beta}} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta}) \\ &= \frac{1}{2\sigma^2} \frac{\partial}{\partial \vec{\beta}} \left(-\vec{X}^T \vec{y} + \vec{X}^T X \vec{\beta} \right) \end{aligned}$$

$$\begin{aligned} \frac{\partial \log L(\vec{\beta})}{\partial \vec{\beta}} = 0 &\implies -\vec{X}^T \vec{y} + \vec{X}^T X \vec{\beta} = 0 \\ &\implies \vec{X}^T X \vec{\beta} = \vec{X}^T \vec{y} \\ &\implies \vec{\beta} = (\vec{X}^T X)^{-1} \vec{X}^T \vec{y} \end{aligned}$$

ML estimate of σ

Maximum Likelihood Estimation: choose σ to maximize $L(\vec{\beta}, \sigma)$ or equivalently $\log L(\vec{\beta}, \sigma)$.

$$\log L(\vec{\beta}, \sigma) = n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{2\sigma^2} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta})$$

ML estimate of σ

Maximum Likelihood Estimation: choose σ to maximize $L(\vec{\beta}, \sigma)$ or equivalently $\log L(\vec{\beta}, \sigma)$.

$$\log L(\vec{\beta}, \sigma) = n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{2\sigma^2} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta})$$

$$\begin{aligned} \frac{\partial \log L(\vec{\beta}, \sigma)}{\partial \sigma} &= -\frac{n}{2} \cdot \frac{1}{2\pi\sigma^2} \cdot 4\pi\sigma - \frac{1}{2} \cdot \frac{-2}{\sigma^3} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta}) \\ &= -\frac{n}{\sigma} + \frac{1}{\sigma^3} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta}) \end{aligned}$$

ML estimate of σ

Maximum Likelihood Estimation: choose σ to maximize $L(\vec{\beta}, \sigma)$ or equivalently $\log L(\vec{\beta}, \sigma)$.

$$\log L(\vec{\beta}, \sigma) = n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{2\sigma^2} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta})$$

$$\begin{aligned} \frac{\partial \log L(\vec{\beta}, \sigma)}{\partial \sigma} &= -\frac{n}{2} \cdot \frac{1}{2\pi\sigma^2} \cdot 4\pi\sigma - \frac{1}{2} \cdot \frac{-2}{\sigma^3} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta}) \\ &= -\frac{n}{\sigma} + \frac{1}{\sigma^3} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta}) \end{aligned}$$

$$\begin{aligned} \frac{\partial \log L(\vec{\beta})}{\partial \vec{\beta}} = 0 &\implies -\frac{n}{\sigma} + \frac{1}{\sigma^3} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta}) = 0 \\ &\implies \sigma^2 = \frac{1}{n} (\vec{y} - X\vec{\beta})^T (\vec{y} - X\vec{\beta}) = \frac{1}{n} \|\vec{y} - X\vec{\beta}\|^2 \end{aligned}$$

Making predictions

Given ML estimates $\vec{\beta}, \sigma^2$, for a new input \vec{x}^* ,

$$p(y_i | \vec{x}^*; \beta, \sigma^2) = N(\vec{\beta}^T \vec{x}^*, \sigma^2)$$

The expected value of y_i is

$$E(y_i) = \vec{\beta}^T \vec{x}^*$$

and 95% confidence interval is given by

$$(\vec{\beta}^T \vec{x}^* - 2\sigma, \vec{\beta}^T \vec{x}^* + 2\sigma)$$

that is,

$$p(\vec{\beta}^T \vec{x}^* - 2\sigma \leq y_i \leq \vec{\beta}^T \vec{x}^* + 2\sigma) = 0.95$$