INFO 6105
Data Science Engineering Methods and Tools

Lecture 6
Lab: Pima Indians Diabetes Data Set

Ebrahim Nasrabadi
nasrabadi@northeastern.edu

College of Engineering
Northeastern University

Fall 2019

R or Python?

Both languages are simple to learn and great for data science, but one might be better than the other depending upon your skills and your priorities.

- experience programming in other languages
- academia or industry
- data processing (pandas, data.table)
- data visualization (matplotlib, ggplot)
- machine learning (sciekit-learn, CARET)
- ...

Scikit-learn

Scikit-learn is a machine learning library for the Python programming language. It provides simple and efficient tools for

- Preprocessing: Feature extraction and normalization.
- Regression: Predicting a quantitative variable
- Classification: Predicting a qualitative variable
- Clustering: Automatic grouping of similar objects into sets.
- Model selection: Comparing, validating and choosing parameters and models.
- Dimensionality reduction: Reducing the number of random variables to consider.

How do I install scikit-learn?

- Option 1: Install scikit-learn library and dependencies (NumPy and SciPy)
- Option 2: Install Anaconda distribution of Python

Scikit-learn

Scikit-learn is a machine learning library for the Python programming language. It provides simple and efficient tools for

- Preprocessing: Feature extraction and normalization.
- Regression: Predicting a quantitative variable
- Classification: Predicting a qualitative variable
- Clustering: Automatic grouping of similar objects into sets.
- Model selection: Comparing, validating and choosing parameters and models.
- Dimensionality reduction: Reducing the number of random variables to consider.

How do I install scikit-learn?

- Option 1: Install scikit-learn library and dependencies (NumPy and SciPy)
- Option 2: Install Anaconda distribution of Python

scikit-learn

Requirements for working with data in scikit-learn

- Features and response should be separate objects
- Features and response should be numeric
- Features and response should be NumPy arrays
- Features and response should have specific shapes

Pima Indians Diabetes Data Set

- Sources: National Institute of Diabetes and Digestive and Kidney Diseases
- Objective: Predict based on diagnostic measurements whether a patient has diabetes
- Number of Instances: 768 (all females and at least 21 years old)
- Number of Features: 8 plus class
  1. Number of times pregnant
  2. Plasma glucose concentration a 2 hours in an oral glucose tolerance test
  3. Diastolic blood pressure (mm Hg)
  4. Triceps skin fold thickness (mm)
  5. 2-Hour serum insulin (mu U/ml)
  6. Body mass index (weight in kg/(height in m)^2)
  7. Diabetes pedigree function
  8. Age (years)
  9. Class variable (0 or 1)
- Download dataset: Machine Learning Repository

Descriptive Statistics

	preg	plas	pres	skin	test	mass	pedi	age	class
count	768.000	768.000	768.000	768.000	768.000	768.000	768.000	768.000	768.000
mean	3.845	120.895	69.105	20.536	79.799	31.993	0.472	33.241	0.349
std	3.370	31.973	19.356	15.952	115.244	7.884	0.331	11.760	0.477
min	0.000	0.000	0.000	0.000	0.000	0.000	0.078	21.000	0.000
25%	1.000	99.000	62.000	0.000	0.000	27.300	0.244	24.000	0.000
50%	3.000	117.000	72.000	23.000	30.500	32.000	0.372	29.000	0.000
75%	6.000	140.250	80.000	32.000	127.250	36.600	0.626	41.000	1.000
max	17.000	199.000	122.000	99.000	846.000	67.100	2.420	81.000	1.000

Observations

There are columns that have a minimum value of zero. On some columns, a value of zero does not make sense and indicates an invalid or missing value.

- Plasma glucose concentration: 5
- Diastolic blood pressure: 35
- Triceps skinfold thickness: 227
- 2-Hour serum insulin: 374
- Body mass index: 11

Descriptive Statistics

	preg	plas	pres	skin	test	mass	pedi	age	class
count	768.000	768.000	768.000	768.000	768.000	768.000	768.000	768.000	768.000
mean	3.845	120.895	69.105	20.536	79.799	31.993	0.472	33.241	0.349
std	3.370	31.973	19.356	15.952	115.244	7.884	0.331	11.760	0.477
min	0.000	0.000	0.000	0.000	0.000	0.000	0.078	21.000	0.000
25%	1.000	99.000	62.000	0.000	0.000	27.300	0.244	24.000	0.000
50%	3.000	117.000	72.000	23.000	30.500	32.000	0.372	29.000	0.000
75%	6.000	140.250	80.000	32.000	127.250	36.600	0.626	41.000	1.000
max	17.000	199.000	122.000	99.000	846.000	67.100	2.420	81.000	1.000

Observations

There are columns that have a minimum value of zero. On some columns, a value of zero does not make sense and indicates an invalid or missing value.

- Plasma glucose concentration: 5
- Diastolic blood pressure: 35
- Triceps skinfold thickness: 227
- 2-Hour serum insulin: 374
- Body mass index: 11

Observations (Cont.)

There are columns that have a minimum value of zero. On some columns, a value of zero does not make sense and indicates an invalid or missing value.

- We can see that columns 1,2 and 5 have just a few zero values, whereas columns 3 and 4 show a lot more, nearly half of the rows.
- This highlights that different "missing value" strategies may be needed for different columns, e.g. to ensure that there are still a sufficient number of records left to train a predictive model.

Missing Values

Why is it important?

- Wrong conclusions
- Some statistical models fail if there is missing values
- Leading to poor performance

How to handle missing values?

- Remove Rows With Missing Values
- Impute Missing Values (e.g., Use a model to replace missing values.)
  - A constant value that has meaning within the domain, such as 0, distinct from all other values.
  - A value from another randomly selected record.
  - A mean, median or mode value for the column.
  - A value estimated by another predictive model.

Missing Values

Why is it important?

- Wrong conclusions
- Some statistical models fail if there is missing values
- Leading to poor performance

How to handle missing values?

- Remove Rows With Missing Values
- Impute Missing Values (e.g., Use a model to replace missing values.)
  - A constant value that has meaning within the domain, such as 0, distinct from all other values.
  - A value from another randomly selected record.
  - A mean, median or mode value for the column.
  - A value estimated by another predictive model.

Lab

Link: Pima Indians Diabetes Data in Python